


Architecture Over Intelligence

How I made AI workflows reliable, even on a small model

Amar Patel · The AI Fellowship Madrid · 9 October 2026



Make one button blue. Nothing else.



Task: make Add to Cart blue 

Opusfived, a parody by Milos Novovic:
opusfived.dev

It changes both buttons and says “Done”.
Every round of clean-up is paid for.



IN SHORT

A clever AI can do more than you asked, call it done, and bill you for the clean-up.

I hit my limit, halfway through a meeting

BEFORE · UNTIL MAY 2026

Claude (cloud)

- Very capable
- Usage limit on my plan
- About €190 of top-ups in six weeks, more than my annual plan

AFTER · FROM MAY 2026

Qwen 3.6 (open model, on my desk)

- Runs on a 16 GB gaming graphics card
- No limit, no bill
- Private: nothing leaves the house

“Economics started the journey.
Drift is what I found along the way.”

IN SHORT

I moved to a smaller, local AI to save money. Then it started to go wrong.

Claude hid the gaps. Qwen didn't.

drift = the AI slowly goes off track, and does not notice

RUN 1 · CLAUDE HAIKU AND SONNET

- ✓ One note, the right date, **5 of 5** quotes to the right person
- ✗ Deleted the transcript **without asking**
- ✗ Haiku made up **6 due dates**

RUN 2 · QWEN, 4-BIT, ON MY OWN COMPUTER

- ✗ Wrote the note **before** it knew who was speaking
- ✗ Named **4 of 4** speakers wrong. The CEO became “head of alliances”
- ✗ Deleted the transcript **without asking**

RUN 3 · SAME QWEN, 8-BIT, IN THE CLOUD

- ✗ **Two** notes, and the **wrong date**
- ✗ **2 of 5** quotes to the wrong person
- ✗ A Chinese word in an English sentence

Same May workflow, same answers, one run each. Test meeting: “OKR Planning with GitLab Executive Team”, GitLab Unfiltered, public on YouTube. Claude is not the May model.

IN SHORT

Claude covered for a weak workflow. Qwen could not.

More instructions. Then SHOUTING.

	11 MAY Before	8 SEP First fix: more words	TODAY What worked: code
Lines in the prompt instructions to the AI	73	416	283
Warnings in CAPITALS MUST, NEVER, STOP...	0	13	0
Scripts doing the sure things normal code	1	4	8

TWO REAL LINES FROM THE 8 SEPTEMBER PROMPT

THIS IS THE ONLY SOURCE OF TRUTH FOR PROGRESS.


If preconditions fail → STOP. Report gate failure. × 9

An instruction is only a suggestion.

IN SHORT

If you write MUST in capitals, your process is missing a check.

My whole flow, split three ways



Between every row, a script checks: is it this step's turn, and did the last step leave its file? If not, stop.

IN SHORT The model thinks. Scripts do. Files remember.

Give each step to the right worker

Model

Does this need judgement?

Who is speaking

Summary

Actions

Priorities

Tone of the meeting

Script

Is there only one right answer?

Date and time

File name and folder

Copy the transcript

Link to the project

Delete old files

5,000 words

In May, the model retyped the whole transcript. It ran out of space, and I paid for every word.

Wrong date

The date was already in the file name. Qwen still got it wrong, in Run 3.

1 → 8 scripts

Today, scripts do the sure things. They get them right every time.

IN SHORT

If the answer never changes, don't ask a model.

Every step leaves a file



A chat forgets

It gets long. Early details get lost. You cannot check it.

4

auto-compactions

X

did not finish

My test, 64K context: after each compaction, the model forgot work it had already finished.

A file remembers

The next step starts from the file, not the chat. A script can check it: **is it there? does it have the right shape?**

0

auto-compactions

✓

finished every step


Same model, same 64K: I cleared the context after every step, on purpose. Each step started again from the files.

IN SHORT

Each step writes down its result, so you can clear the context between steps.

Keep progress outside the model


● progress file .process-state.json	
current_step	2.6
completed_steps	-1, 0, 1, 2, 2.5, 2.6
speakers confirmed	yes
aborted	no



In rehearsal I closed everything in the middle of finding actions. When I opened it again, it read this file and carried on.
Nothing repeated. Nothing lost.

IN SHORT The model forgets everything. The file doesn't.

Check before and after every step



WHAT IT CAUGHT

✓ SPEAKERS

The model still guessed **2 of 4** wrong. But it had to **stop and show me first**. One-line fix, before anything was written.

✓ MY OWN BUG

A time-zone bug looked for a **12:00** note. The meeting was at **10:00**. It **refused to file** instead of filing wrong.

✓ A SKIPPED REVIEW

A review is done only **after I reply**. The model tried to **skip my actions review**. The check stopped it.

IN SHORT

The check is a small script, not a sentence in the prompt. The model cannot talk it round.

Mark the one-way doors

WHO DECIDES? ME, WHEN A STEP IS...

Ambiguous Name speakers

Sensitive Draft follow-up

Irreversible Clean up

Confirm

Yes or no

One-way door

One-way door · last step: clean up

Delete the transcript, audio and progress file for this meeting. Approve?

Approve

Block

Only I can click. After Block, the shell locks down.

Retests: 3 of 3 stopped and asked.

IN SHORT

A quick look where it's unclear. A locked door where it can't be undone.

My flow in May vs my flow today

SAME MEETING, MODEL AND ANSWERS	MAY		TODAY	
Checked the speakers before writing				
Quotes given to the right person		3 of 5		5 of 5
One note per meeting		2 notes		1 note
Right date				
Asked before deleting				
Carried on after I quit		not tested		

To be fair: today's flow also has more steps and better scripts. This is not a lab test. It is my flow in May against my flow today.

IN SHORT The old flow went wrong quietly. The new one goes wrong loudly, early, and where I can fix it.

Smaller models, bigger jobs

ROOM TO BREATHE

A smaller model does the judgement. Scripts do the rest.

No tokens spent on work with one right answer. Meetings stop using the capacity I need for everything else.

SAFER

Every step leaves a file. Every Approve and Block is logged.

You can show exactly what happened, and who said yes.

PORTABLE

Swap the model without rewriting the flow.

Less lock-in. Data can stay on your machine, if a model that fits clears the bar.

A SMALLER MODEL. NOT THE SMALLEST · ONE MEETING, FIVE MODELS


One couldn't read the steps. Same Qwen at **3-bit** invented a speaker who was in Vegas. At **8-bit** it got everyone right. **Architecture lowers the bar. It doesn't remove it.**

The same shape fits: sales call follow-ups · invoice processing · customer onboarding · hiring · month-end close

IN SHORT

Reliable flows let you use smaller models on bigger jobs.

Models are splitting: some write, some decide



WHY I DON'T PICK A WINNER · JEV DECISION INDEX 0.3, 7 OCT

Three weeks after Jev's early access opened, a community index ranked **112 deciders**. Jev came **third**.

IN SHORT

Whoever makes the best decider, it slots into one box. The flow stays the same.

Five rules. Start with a checklist.

- 1 Write the steps down,** and give each one to the right worker. Judgement → model. One right answer → script.
- 2 Make every step leave a file.** The next step starts from the file, not the chat.
- 3 Keep progress outside the model.** A note says which step you are on.
- 4 Check before and after every step.** If the answer is no, stop.
- 5 Mark the one-way doors.** A person decides what is ambiguous, irreversible or sensitive.

Try one tonight: take your messiest AI workflow. Write down the steps, and what each one should produce.

IN SHORT

Without code, you are the check. Add scripts where the answer never changes, and test them.

Models change every month.

“ Build your flow so the model is a part you can swap.



Connect

[linkedin.com/in/amarapatel](https://www.linkedin.com/in/amarapatel)



The talk

amarapatel74.github.io

Thank you
Amar Patel

Links from the talk

- **Opusfived**, the "make one button blue" game: opusfived.dev
- **Jev Decision Index**, the ranking on slide 14: huggingface.co/spaces/multimodalart/jev-decision-index
- **Jev launch post**, TypeSafe AI: typesafe.ai/blog/introducing-system-one-models-and-jev
- **"The State Machine Is the Agent"**, code owns the loop, the model judges at the branches: stacktoheap.com/blog/2026/09/21/the-state-machine-is-the-agent
- **Temporal**, the same idea at company scale: temporal.io